

COMP 2402

Random Binary Search Trees


Alina Shaikhet

Let's build a Tree

We want to create a binary search tree containing the integers **0,...,14**.

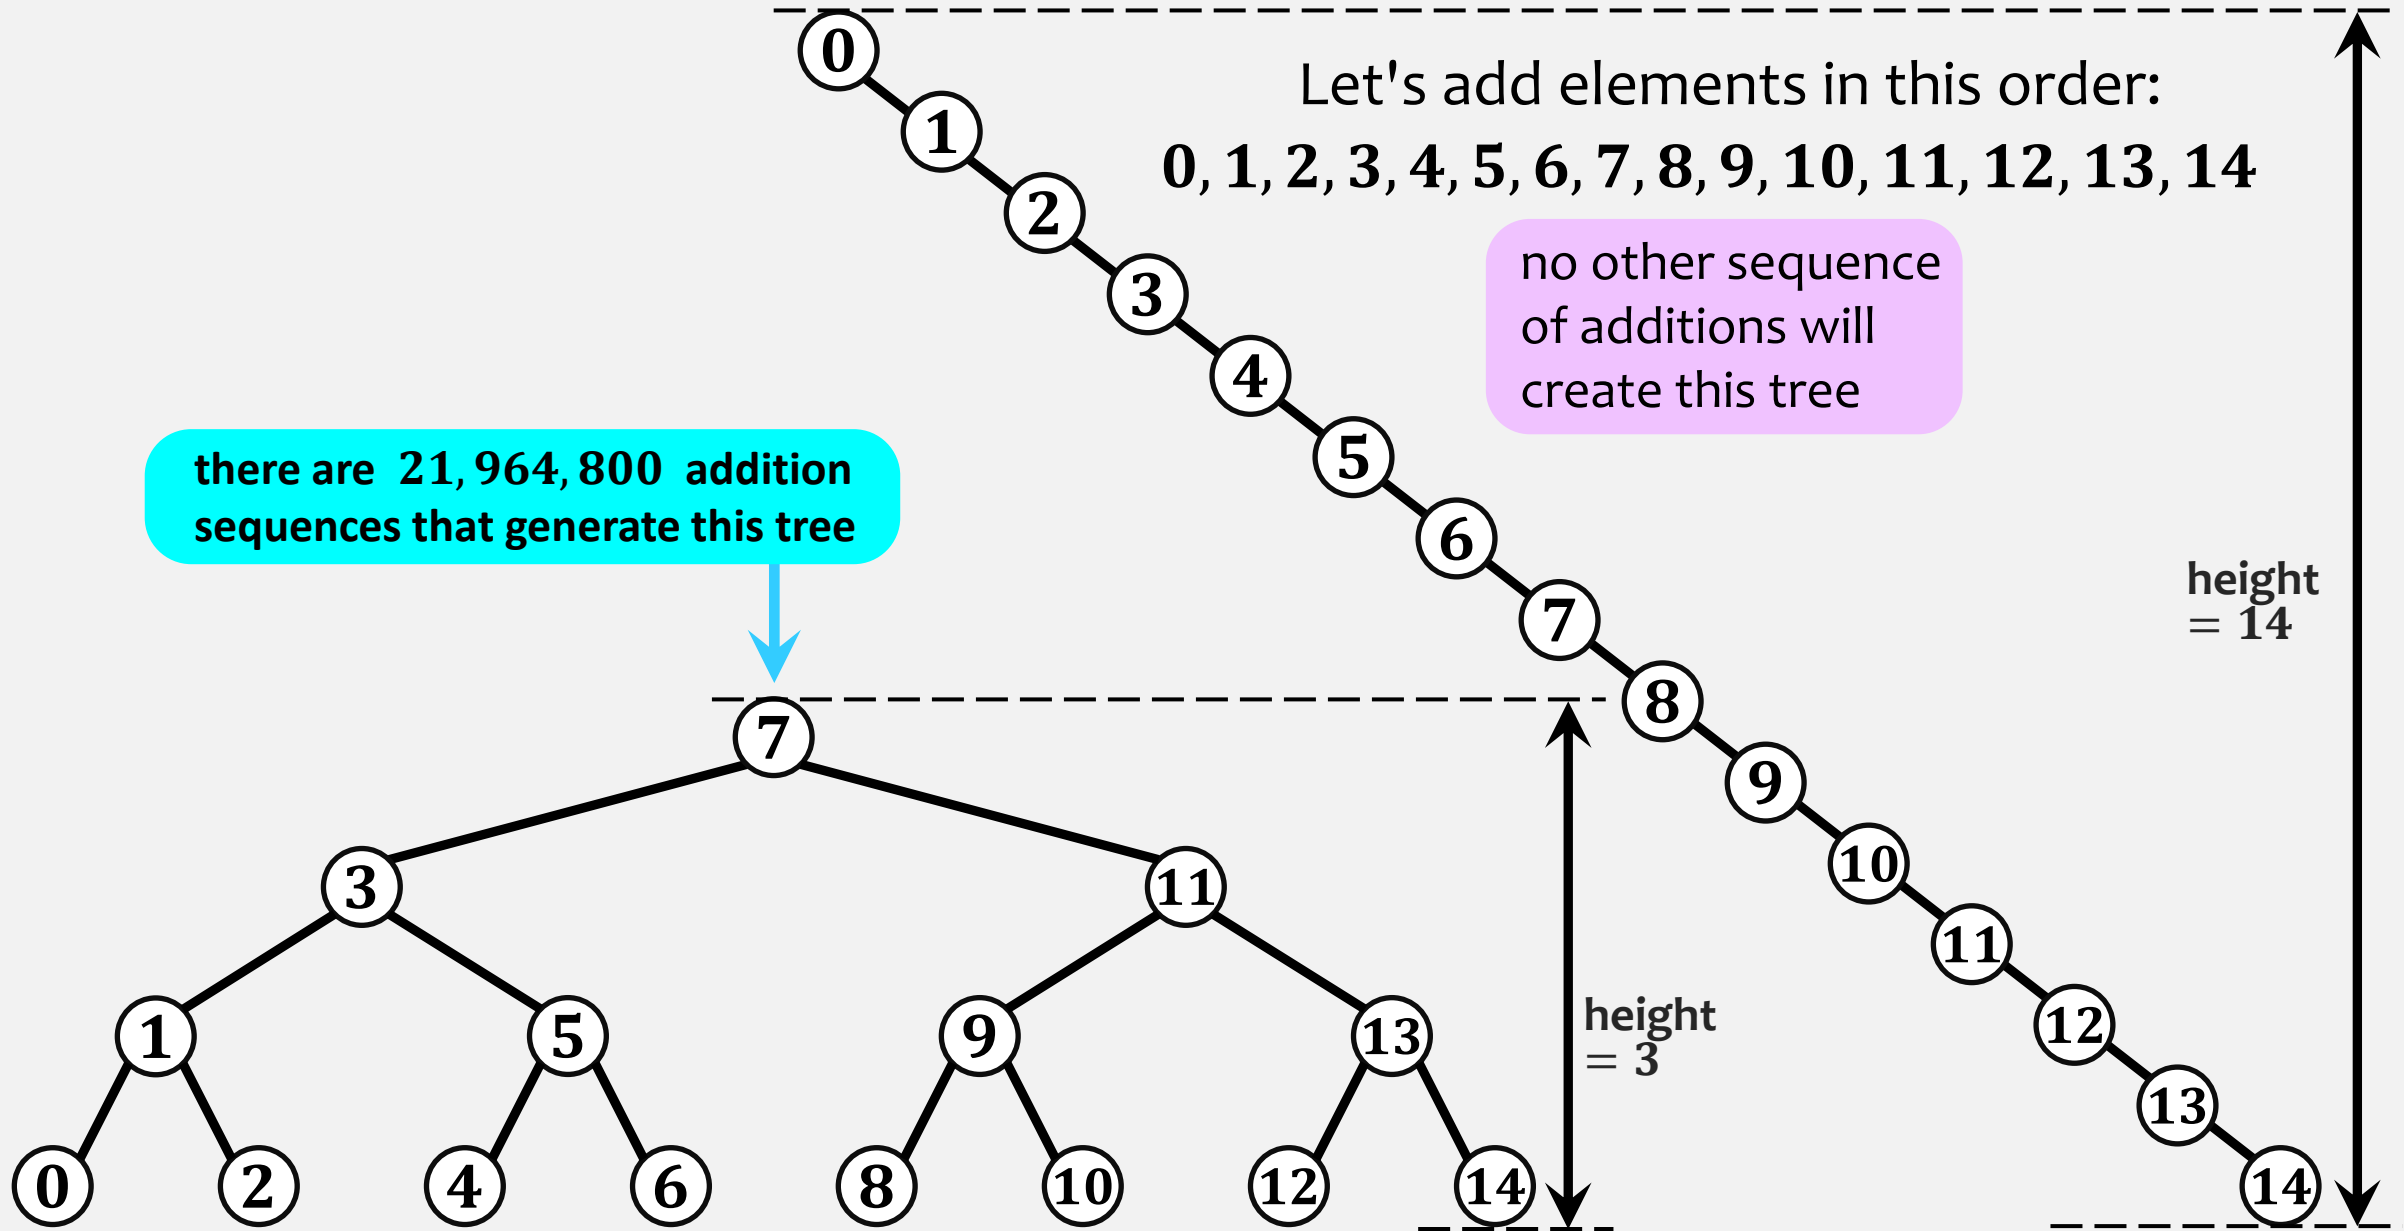
- Start with an empty tree,
- Add elements to the tree one at a time.

Shape of the tree will depend on the order, in which we add elements to the tree.



Let's add elements in this order:
7, 3, 5, 11, 13, 1, 4, 9, 12, 2, 6, 14, 0, 10, 8
or this:
7, 11, 9, 3, 1, 2, 8, 5, 10, 4, 13, 0, 12, 6, 14

Two binary search trees containing the integers 0, ..., 14



Random Binary Search Tree

- Take a **random permutation** of n distinct values (ranks) (for example $0, \dots, n - 1$).
- Insert elements according to their ranks one at a time into a binary search tree.

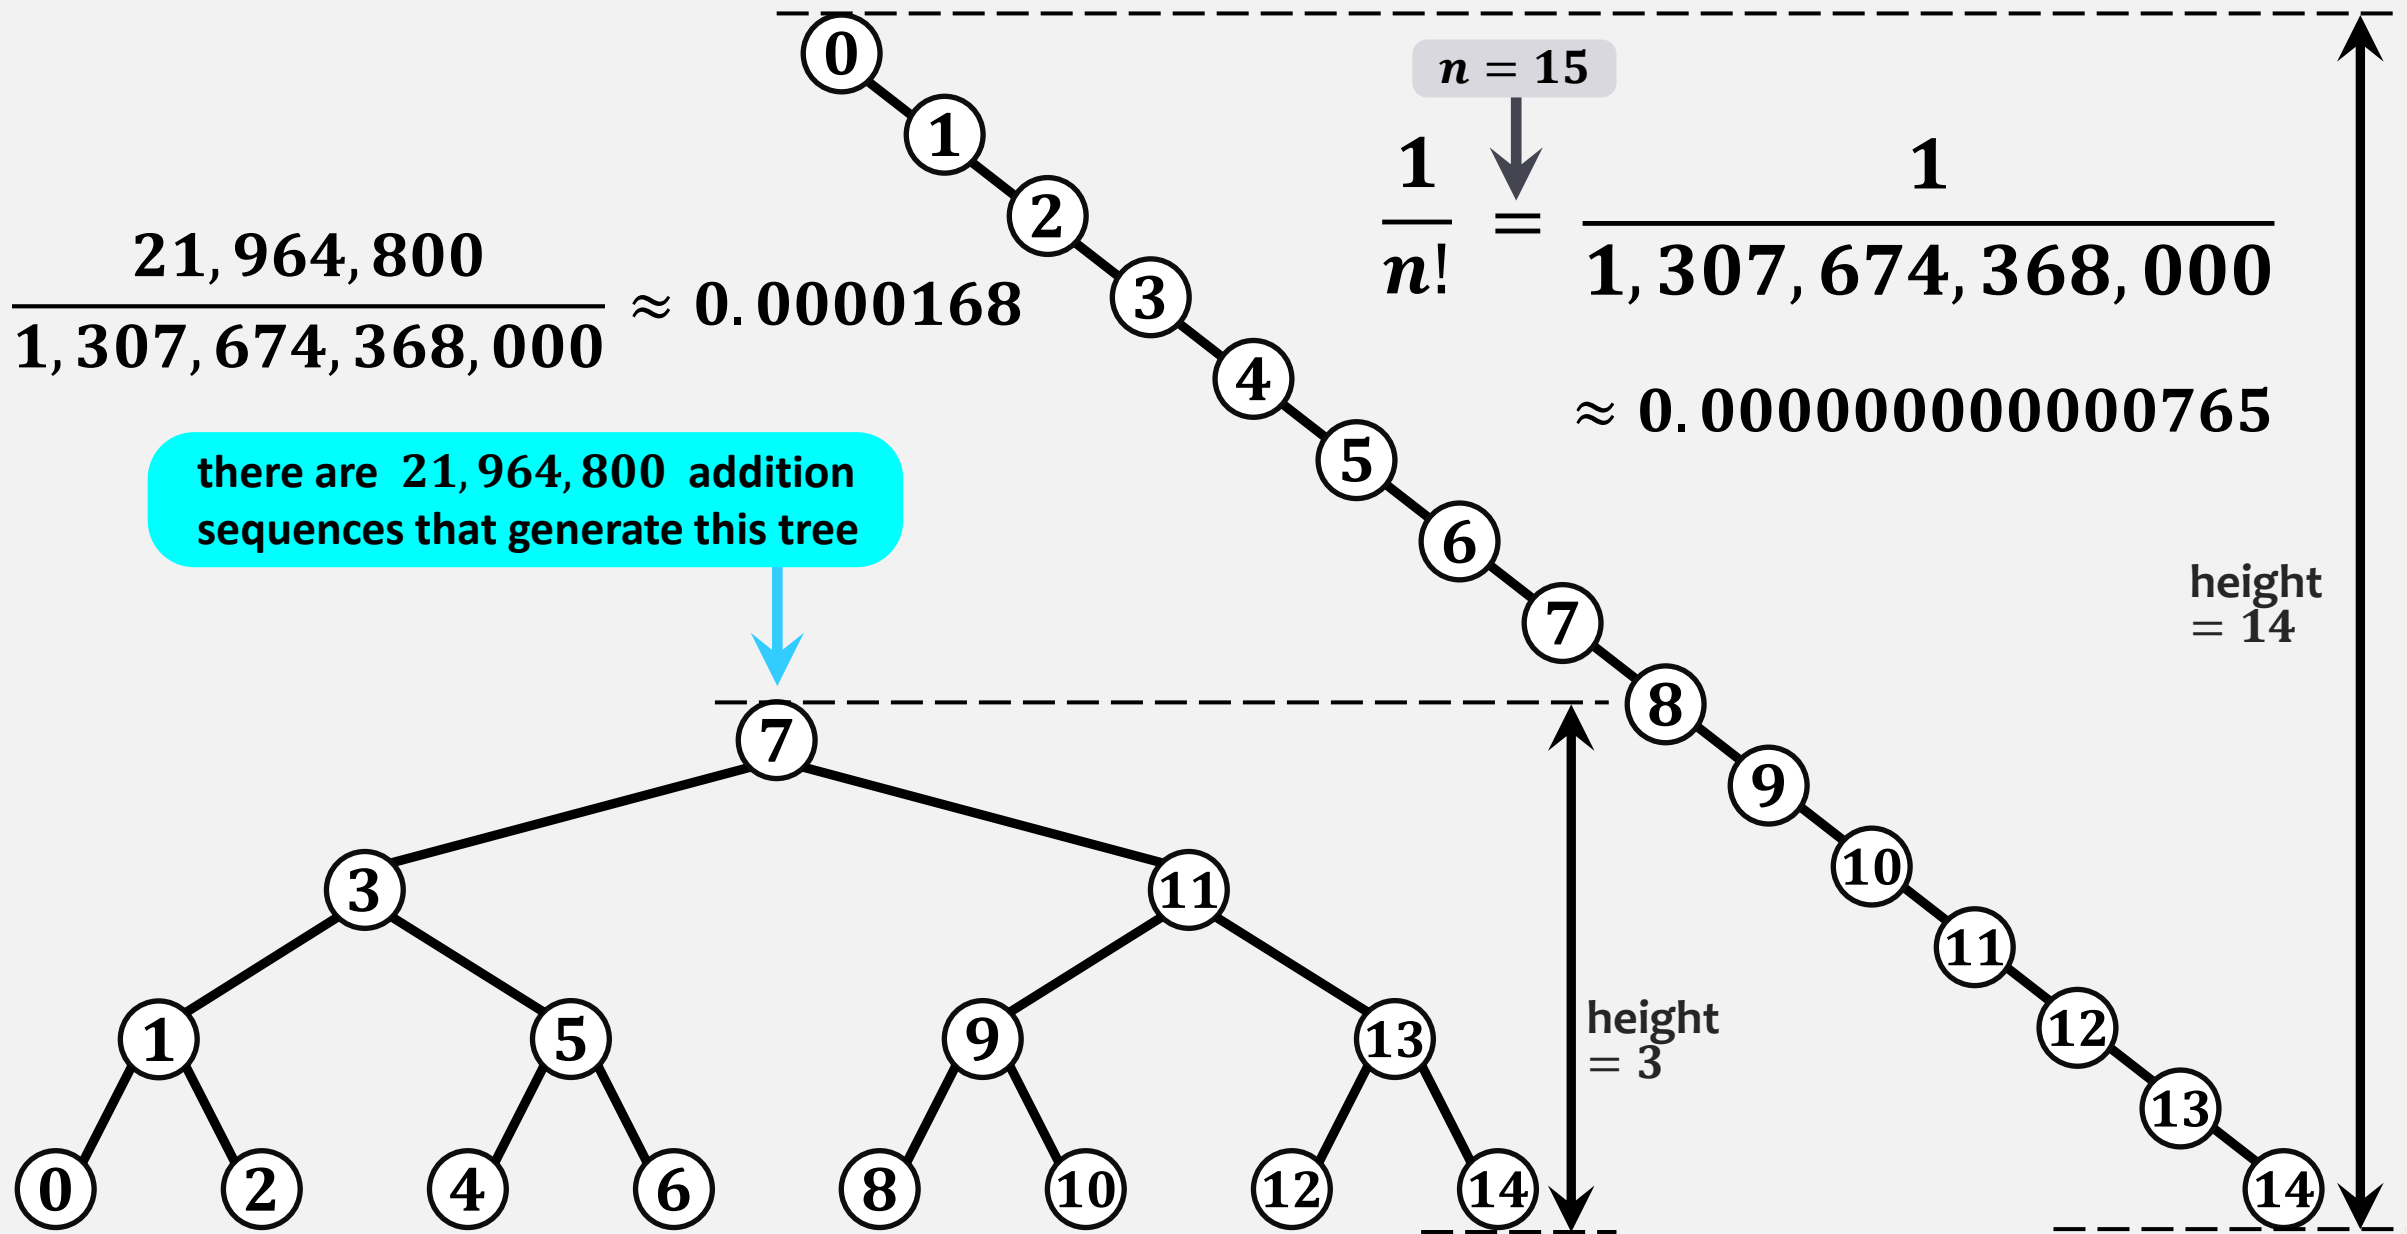
In how many ways can we permute n elements?

$$n! = 1 \times 2 \times 3 \times \dots \times (n - 1) \times n$$

By **random permutation** we mean that each of the possible $n!$ permutations (orderings) of $0, \dots, n - 1$ is equally likely,

So, the probability of obtaining any particular permutation is $\frac{1}{n!}$

Two binary search trees containing the integers 0, ..., 14



Lemma 7.1

In a random binary search tree with n nodes, the expected length of the **search path** for any value x (whether x is stored in the tree or not) is **at most**


$$2 \ln n + 2 \approx 1.38 \log_2 n + 2$$

Recall: The expected length of a search path in a SkipList is at most $2 \log_2 n + O(1)$.

This result is with respect to the **random permutation** used to create the random binary search tree. We cannot maintain this result under $\text{add}(x)$ and $\text{remove}(x)$ operations.

The problem with random binary search trees is that **they are not dynamic**.

Random binary search trees don't support the $\text{add}(x)$ or $\text{remove}(x)$ operations.

So, we cannot use them to implement the **SSet** interface.


Heap

Heap is a binary tree where each node u has a **priority** $u.p$

(Min-)Heap Property:

For each node u except the root:

$$u.\text{parent}.p < u.p$$




The node with **minimum priority** has to be the **root**.

You cannot effectively search for things in heaps. They are not designed for this.

Treap

A **Treap** is a binary tree T where each node has:

- a value x (**key**)
- a priority p



T is a **binary search tree** with respect to the x values:

$$x_1 < x < x_2$$

T is a **heap** with respect to the p values. (each node has a priority smaller than that of its two children):

$$p < p_1 \text{ and } p < p_2$$

Priority values in a Treap are **unique** and **assigned at random**.

Treap

x	p
-----	-----



- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$.
All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively **root**

11	5	15	3	18	7	20	8	21	1	25	9	29	2
----	---	----	---	----	---	----	---	----	---	----	---	----	---

Treap

x	p
-----	-----


- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively



Treap

x	p
-----	-----

- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively



$\underbrace{11 \mid 5}$
 $x < 15$



$\underbrace{18 \mid 7 \mid 20 \mid 8}$
 $15 < x < 21$

$\underbrace{25 \mid 9 \mid 29 \mid 2}$
 $x > 21$

Treap

x	p
-----	-----


- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively



Treap

x	p
-----	-----



- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively



Treap

x	p
-----	-----



- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively



Treap

x	p
-----	-----


- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively




Treap

x	p
-----	-----


- Node with the smallest priority is a root (r).
- All nodes with keys **smaller** than $r.x$ are stored in the subtree rooted at $r.left$. All nodes with keys **larger** than $r.x$ are stored in the subtree rooted at $r.right$.
- Repeat recursively




Treap




Treap



sorted by priority →




Treap



sorted by priority →




Treap



sorted by priority →




Treap




sorted by priority →




Treap



Treap





Treap



Treap

Alternative definition of a Treap:

We can think of a Treap as a BinarySearchTree whose nodes were added in **increasing order of priority**.



Treap, where the **priorities** are **assigned at random**, is a **Random Binary Search Tree**

Lemma 7.2

In a Treap (where the **priorities** are **assigned at random**) with n nodes, the expected length of the **search path** for any value x (whether x is stored in the treap or not) is **at most**

$$2 \ln n + 2 \approx 1.38 \log_2 n + 2$$

Treap: SSet implementation

- $\text{add}(x)$,
- $\text{remove}(x)$,
- $\text{find}(x)$ – find the **smallest** value that is $\geq x$.



Search in a Treap for x . It will take $O(\log n)$ **expected** time.

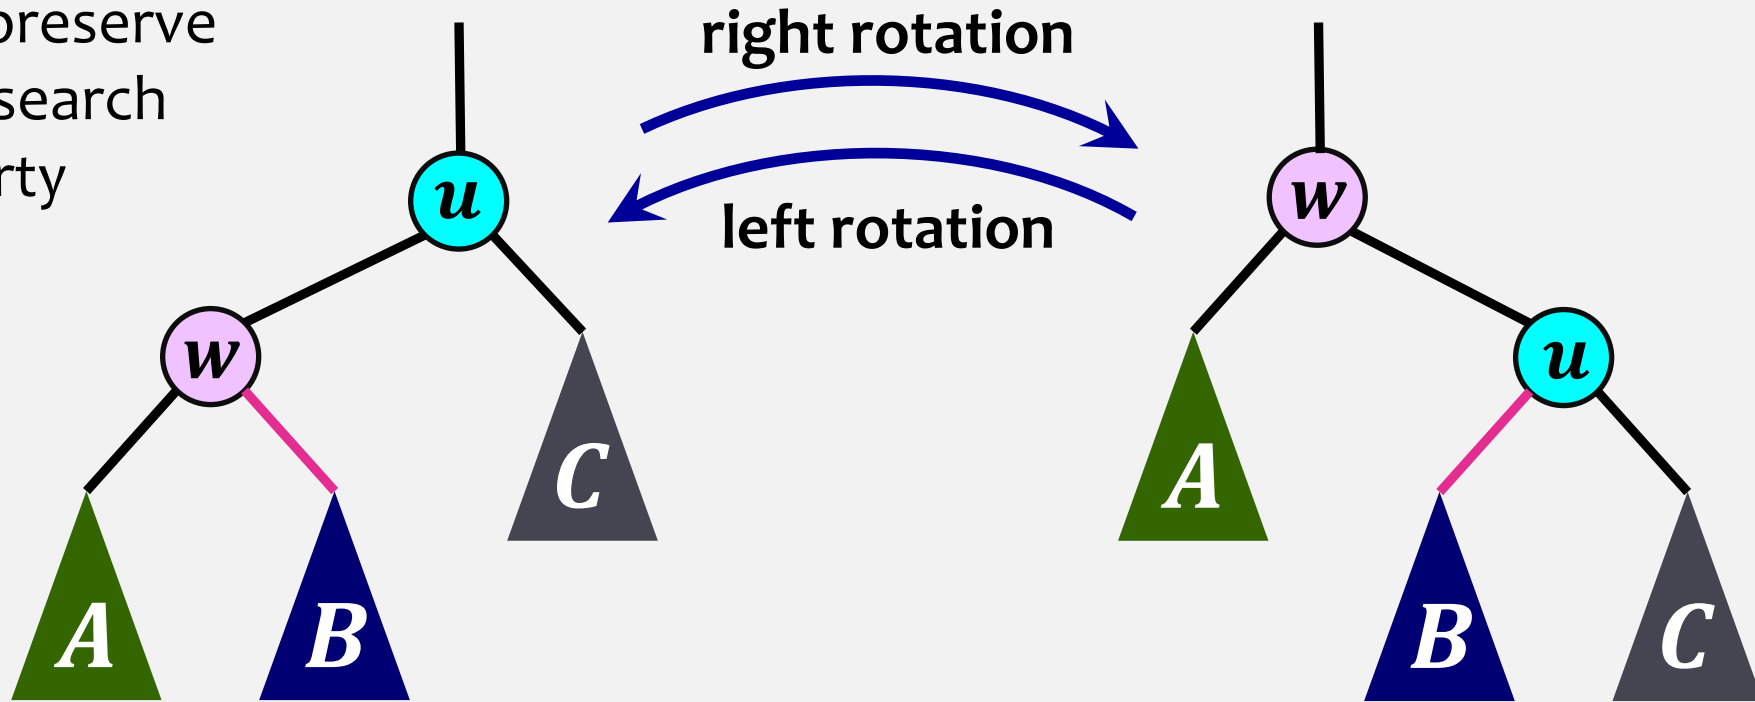
if we can maintain a random treap

To support the $\text{add}(x)$ and $\text{remove}(x)$ operations, a treap needs to perform **rotations** in order to maintain the heap property.

Rotation

the most important property of a rotation (right) is that the depth of w decreases by one while the depth of u increases by one.

Rotations preserve the binary search tree property




$$A < w < B < u < C$$

Any of the trees A , B or C can be empty.

Treap: SSet – add(x)



- Search for x in the tree
- If it is not there, add x as a leaf: q



add(**28**)

Treap: SSet – add(x)



- Search for x in the tree
- If it is not there, add x as a leaf: q
- Give it a random priority: $q.p$
- As long $q.p < q.\text{parent}.p$ rotate it.



add(28)

Treap: SSet – add(x)



- Search for x in the tree
- If it is not there, add x as a leaf: q
- Give it a random priority: $q.p$
- As long $q.p < q.\text{parent}.p$ rotate it.



add(28)

Treap: SSet – add(x)

- Search for x in the tree
- If it is not there, add x as a leaf: q
- Give it a random priority: $q.p$
- As long $q.p < q.\text{parent}.p$ rotate it.



add(28)

Treap: SSet – add(x)

- Search for x in the tree
- If it is not there, add x as a leaf: q
- Give it a random priority: $q.p$
- As long $q.p < q.\text{parent}.p$ rotate it.

$O(\log n)$


$O(1)$

$O(1)$

$O(\log n)$

What is the running time of add(x) operation?

$O(\log n)$ expected




Each time we rotate we get closer to the root. That means that we can't rotate more times than the length of the search path.

Each rotation takes $O(1)$ time (it affects at most 6 pointers).

Recall – Binary Search Trees – $\text{remove}(x)$

$\text{remove}(2)$




To delete a value stored in **BinarySearchTree**:

- Find node u that contains value x .
- If u is a leaf, then we can just detach u from its parent.

Recall – Binary Search Trees – $\text{remove}(x)$


$\text{remove}(2)$
 $\text{remove}(7)$



To delete a value stored in **BinarySearchTree**:

- Find node u that contains value x .
- If u is a leaf, then we can just detach u from its parent.
- If u has only one child, then we can splice u from the tree by having $u.\text{parent}$ adopt u 's child.

Recall – Binary Search Trees – $\text{remove}(x)$




To delete a value stored in **BinarySearchTree**:

- Find node u that contains value x .
- If u is a leaf, then we can just detach u from its parent.
- If u has only one child, then we can splice u from the tree by having $u.\text{parent}$ adopt u 's child.
- If u has two children, then find a node w , that has less than two children, such that $w.x$ can replace $u.x$.

Choose w , such that $w.x$ is the smallest value in the subtree rooted at $u.\text{right}$. This node has no left child.

Recall – Binary Search Trees – $\text{remove}(x)$





To delete a value stored in **BinarySearchTree**:

- Find node u that contains value x .
- If u is a leaf, then we can just detach u from its parent.
- If u has only one child, then we can splice u from the tree by having $u.\text{parent}$ adopt u 's child.
- If u has two children, then find a node w , that has less than two children, such that $w.x$ can replace $u.x$.

Choose w , such that $w.x$ is the smallest value in the subtree rooted at $u.\text{right}$. This node has no left child.

Treap: SSet – remove(x)

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.



remove(21)



$q.\text{left}.p > q.\text{right}.p$

$3 > 1.5$

do **Left** Rotation

Treap: SSet – remove(x)

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.



Note: we increased the height of the tree

remove(21)



$$q.\text{left}.p < q.\text{right}.p$$

$$3 < 9$$

do **Right** Rotation

Treap: SSet – remove(x)

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.



remove(21)



$q.\text{left}.p < q.\text{right}.p$

$7 < 9$

do **Right** Rotation

Treap: SSet – remove(x)

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.





remove(21)

q .left = null

do **Left** Rotation

Treap: SSet – $\text{remove}(x)$

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.




$\text{remove}(21)$

$q.\text{left} = q.\text{right} = \text{null}$


Remove the leaf 21

Treap: SSet – $\text{remove}(x)$

- Find node q that contains value x .
- If q is a leaf, then we can just detach q from its parent.
- If q is not a leaf, then rotate it down until it is a leaf; then detach.



$\text{remove}(21)$




Treap: SSet – remove(x)

- Find node q that contains value x . $O(\log n)$
- If q is a leaf, then we can just detach q from its parent. $O(1)$
- If q is not a leaf, then rotate it down until it is a leaf; then detach. $O(\log n)$

What is the running time of remove(x) operation?

$O(\log n)$ expected



Think about adding node 21|1

Each time we rotate we get further away from the root. That means that we can't rotate more times than the length of the longest path (which is $n - 1$).

The time it takes to do a deletion on a tree with n nodes is equal to the time it takes to do an insertion on a tree with $n - 1$ nodes.

Theorem 7.2

A **Treap** implements the **SSet** interface. A **Treap** supports the operations $\text{add}(x)$, $\text{remove}(x)$, and $\text{find}(x)$ in $\mathcal{O}(\log n)$ expected time per operation.

The search paths in a **Treap** are considerably shorter compared to **SkipLists** and this translates into noticeably faster operations.