

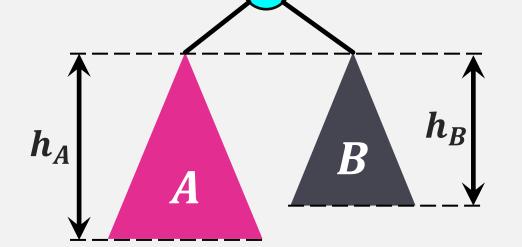
#### **Balanced Search Trees**

A node in a tree is **height-balanced** if the heights of its subtrees differ by no more than 1

If  $|h_A - h_B| \le 1$  then u is **height-balanced.** 

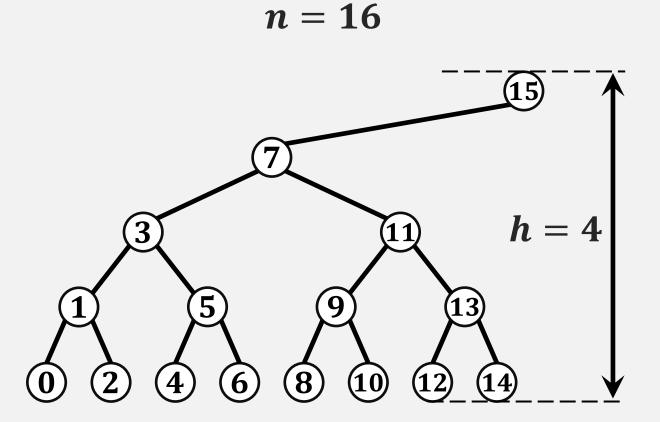
A tree is height-balanced iff every node in it is height-balanced.

When we say that a tree is **balanced**, we usually mean that its height is  $c \log n$ , where c is some constant and n is the number of elements in the tree.



#### **Balanced Search Trees**

Is it possible to make the height of this tree smaller?



A perfect binary tree is a binary tree in which all interior nodes have two children, and all leaves have the same depth (or same level).

### Scapegoat Trees

A **ScapegoatTree** keeps itself balanced by performing partial rebuilding operations (where an entire subtree is deconstructed and rebuilt into a perfectly balanced subtree).

A **ScapegoatTree** is a **BinarySearchTree** that, in addition to keeping track of the number n of nodes in the tree also keeps a counter q that maintains an **upper-bound** on the number of nodes.

$$n \leq q \leq 2n$$

$$q/2 \le n \le q$$

A **ScapegoatTree** can look surprisingly unbalanced, however, it always maintains logarithmic height.

The height of the ScapegoatTree does not exceed

$$\log_{3/2} q \le \log_{3/2} 2n < \log_{3/2} n + 2$$

## Logarithms

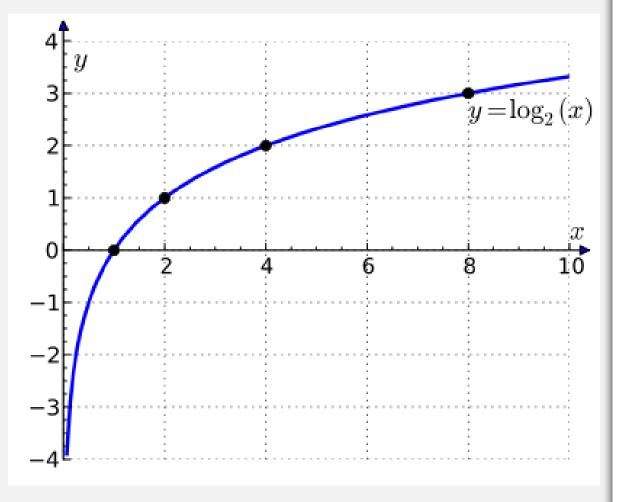
$$y = \log_b n$$

 $\log_b n$  is the number of times we divide n by b before it becomes  $\leq 1$ .

$$b^{y} = b^{\log_b n} = n$$

$$2^8 = 256$$
  
 $8 = \log_2 256$ 

$$\log_b xy = \log_b x + \log_b y$$



$$\log_b n = \frac{\log_a n}{\log_a b} = \frac{\log_a n}{c} = \frac{1}{c} \log_a n$$

That is why the logarithm base is irrelevant when using Big *O* notation.

## SSet Implementation

- add(x),
- remove(*x*),
- $O(\log n)$  ignoring rebuild()
- find(x) find the smallest value that is  $\geq x$ .  $\leftarrow$   $O(\log n)$

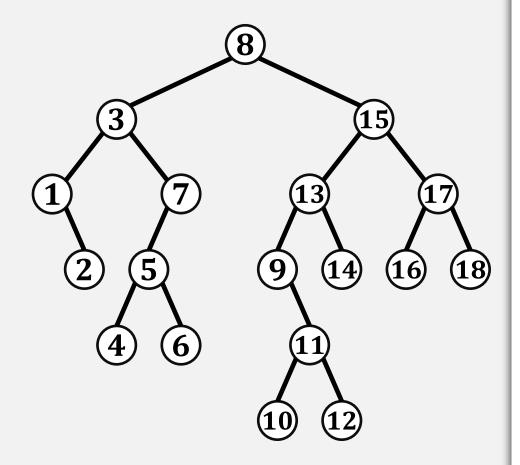
find(x) is implemented using the standard algorithm for searching in a **BinarySearchTree**. This takes time proportional to the height of the tree which is

$$< \log_{3/2} n + 2$$

Use the usual algorithm for deleting the value x stored in **BinarySearchTree**:

- Find node u that contains value x.
- If u is a leaf, then detach u from its parent.
- If u has only one child, then splice u from the tree by having u.parent adopt u's child.
- If u has two children, then find a node w, that has less than two children, such that w.x can replace u.x. Then remove w. Option 1: Find the largest element in the left subtree of u,

Option 2: Find the smallest element in the right subtree of u.

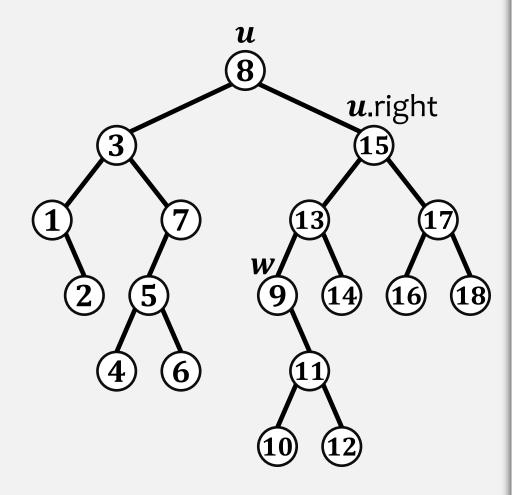


Let's choose option 2.

Start at u.right and keep going to the left child until there is no left child. The node you stopped at is w.

- w. x is the smallest value in the subtree rooted at u.right.
- w has no left child.

remove(8)

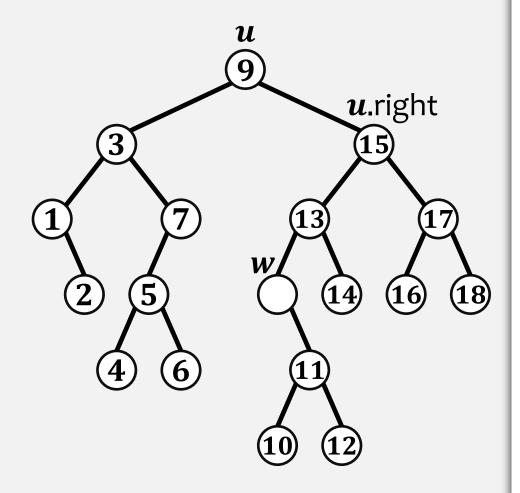


Let's choose option 2.

Start at u.right and keep going to the left child until there is no left child. The node you stopped at is w.

- w. x is the smallest value in the subtree rooted at u.right.
- w has no left child.

remove(8)



delete node w

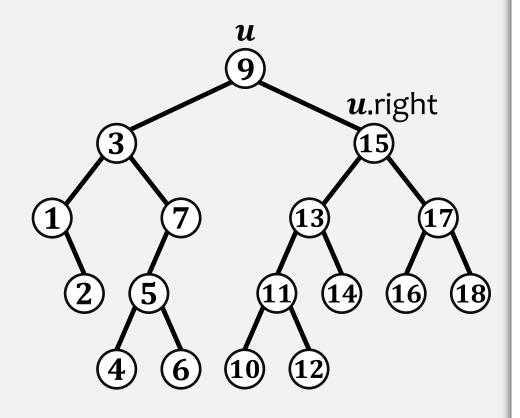
Let's choose option 2.

Start at u.right and keep going to the left child until there is no left child. The node you stopped at is w.

- w. x is the smallest value in the subtree rooted at u.right.
- w has no left child.

Note that remove(x) can never increase the height of the tree.

remove(8)



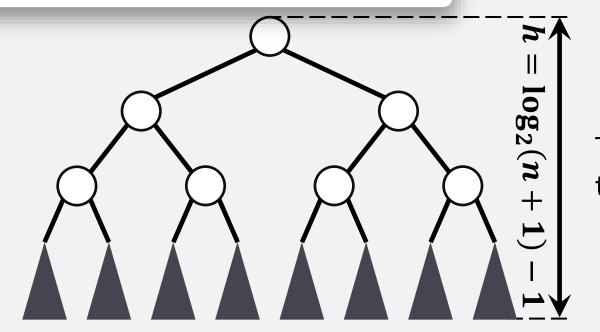
delete node w

At all times:  $n \le q \le 2n$ 

- remove x using standard BST algorithm for removal.
- n -
- if q>2n then rebuild(root) and set q=n

$$2^0 + 2^1 + \dots + 2^k = 2^{k+1} - 1$$

this function rebuilds the tree into a perfectly balanced tree



For Complete Binary Tree

$$h = \lfloor \log_2 n \rfloor$$

There are  $2^{h+1} - 1$  nodes in a full binary tree of height h:

ight 
$$h$$
:  $2^{h+1} - 1 = n$   $2^{h+1} = n + 1$ 

$$h + 1 = \log_2 2^{h+1} = \log_2 (n+1)$$

At all times:  $n \le q \le 2n$ 

- remove x using standard BST algorithm for removal.
- n -
- if q > 2n then rebuild(root) and set q = n 0(n)

this function rebuilds the tree into a perfectly balanced tree

After rebuild() for every node u:

 $|size(u.left) - size(u.right)| \le 1$ 

If we ignore the cost of rebuilding, the running time of the remove(x) operation is proportional to the height of the tree and is therefore  $O(\log n)$ .

# SSet: add(x)

At all times:  $n \le q \le 2n$  $h \le \log_{3/2} q < \log_{3/2} n + 2$ 

- Use the standard algorithm for adding x to **BinarySearchTree**:
- n++
- q++
- If  $(\operatorname{depth}(u) > \log_{3/2} q)$  then walk from uback up to the root looking for a scapegoat w:

$$\frac{\text{size}(w.\text{child})}{\text{size}(w)} > \frac{2}{3}$$
 w. child is a child of w on the path from the root to a

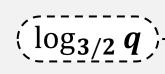
from the root to  $\boldsymbol{u}$ 

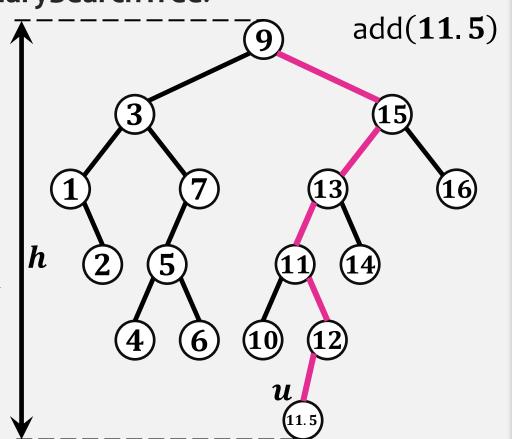
Then rebuild(w) O(size(w))

If we ignore the cost of finding the scapegoat and rebuilding:

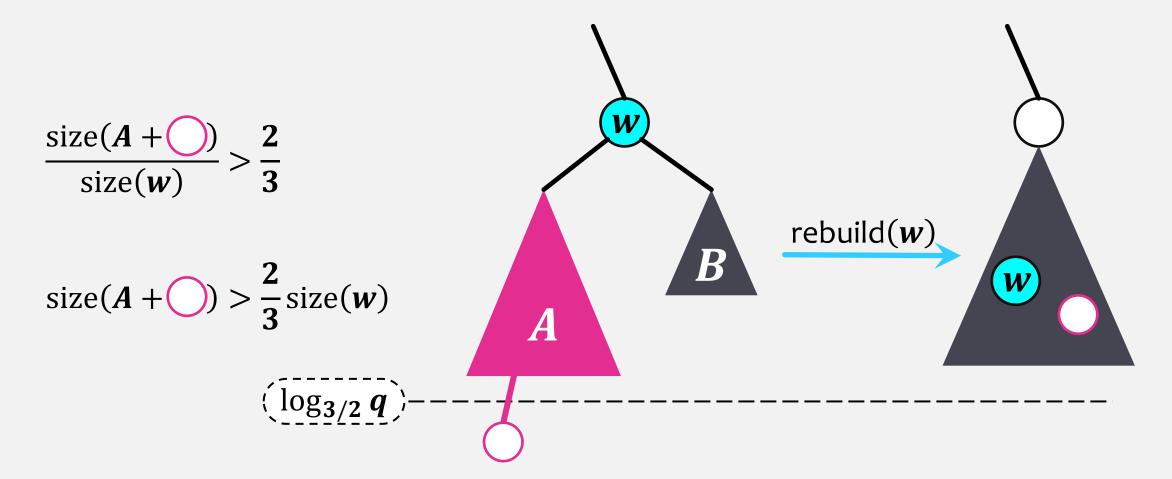
 $O(\log n)$ 

add(x) can increase the height of a tree





## SSet: add(x) – scapegoat w



Does the scapegoat exist?

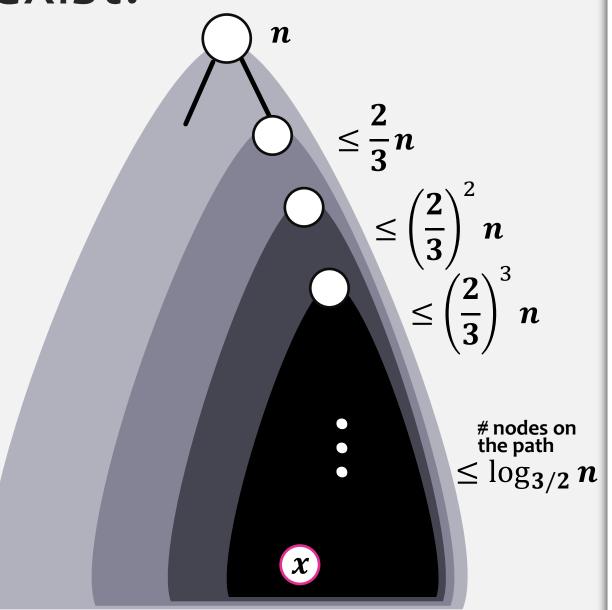
At all times:  $n \le q \le 2n$ 

Suppose, that there does not exist a scapegoat  $\boldsymbol{w}$  on the path from  $\boldsymbol{u}$  (with value  $\boldsymbol{x}$ ) to the root.

Then for all the nodes w on the path

$$\frac{\operatorname{size}(w)}{\operatorname{size}(\operatorname{parent of }w)} \leq \frac{2}{3}$$

Thus, the length of the path from the root to u is at most  $\log_{3/2} n \leq \log_{3/2} q$ So, we didn't cross the  $(\log_{3/2} q)$  line. Contradiction!



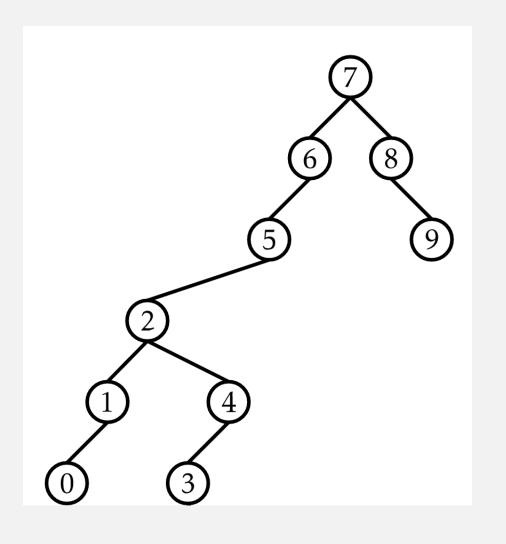
#### Scapegoat Trees

A **ScapegoatTree** can look surprisingly unbalanced, however, it always maintains logarithmic height:  $h \le \log_{3/2} q$ 

For the tree in the example:

$$n = 10 = q$$
  
height = 5

$$5 < \log_{3/2} 10 \approx 5.679$$



## Scapegoat Trees

A **ScapegoatTree** can look surprisingly unbalanced, however, it always maintains logarithmic height:  $h \le \log_{3/2} q$ 

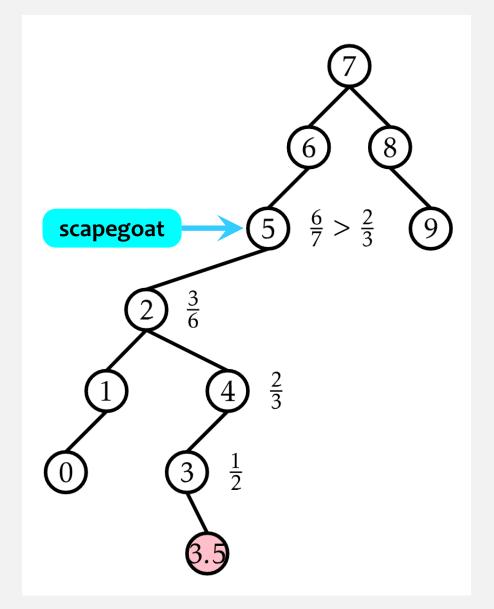
For the tree in the example:

$$n = 10 = q$$

$$height = 56$$

$$5 < \log_{3/2} 10 \approx 5.679$$

$$6 \ge \log_{3/2} 11 \approx 5.914$$



from the ods textbook

# SSet: add(x)

```
boolean add(T x) {
  // first do basic insertion keeping track of depth
  Node < T > u = newNode(x);
  int d = addWithDepth(u);
  if (d > log32(q)) {
    // depth exceeded, find scapegoat
    Node<T> w = u.parent;
    while (3*size(w) <= 2*size(w.parent))</pre>
      w = w.parent;
    rebuild(w.parent);
  return d \ge 0;
```

## rebuild(u)

- traverse (in-order) u's subtree and collect all its nodes into an array a
- recursively build a balanced subtree using a:

$$m=\frac{a.\operatorname{length}}{2}$$

a[m] becomes the root of the new subtree, a[0], ..., a[m-1] get stored recursively in the **left** subtree, and a[m+1], ..., a[a] length a[m+1] get stored recursively in the right subtree.

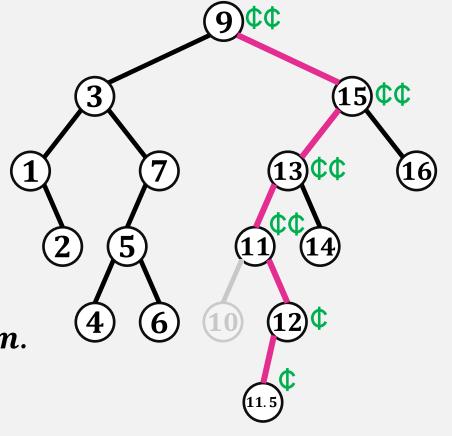
**Lemma 8.3:** Starting with an empty **ScapegoatTree** over any sequence of m add(x) and remove(x) operations the total time spent in calls to rebuild(u) is  $O(m \log N)$ , where N is the maximum value of n at any time.

#### Credit scheme:

- each node stores a number of credits
- during an insertion or deletion of u, we give one credit to each node on the path to u.
- during a deletion we also store an additional credit

#### Invariant:

- 1. The number of q-credits saved up is at least q n.
- 2. The number of credits on any node u is at least  $|\operatorname{size}(u, \operatorname{left}) \operatorname{size}(u, \operatorname{right})| 1$



Invariant:

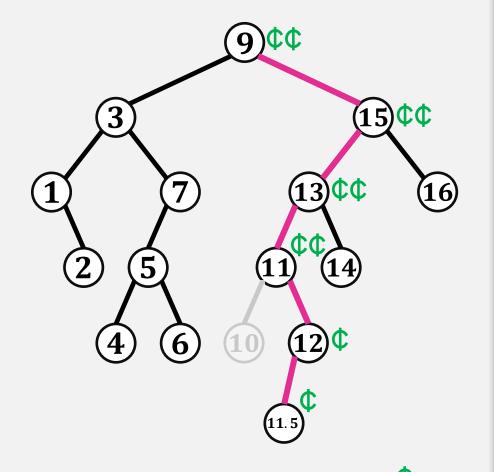
- 1. The number of q-credits saved up is at least q-n.
- 2. The number of credits on any node u is at least  $|\operatorname{size}(u, \operatorname{left}) \operatorname{size}(u, \operatorname{right})| 1$

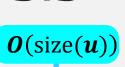
If we call rebuild(u) during a deletion, it is because we decremented n and now q>2n.

But we have q - n q-credits stored.

$$q-n>2n-n=n$$

We use these to pay for the O(n) time it takes to rebuild the root.







- 1. The number of q-credits saved up is at least q-n.
- 2. The number of credits on any node u is at least  $|\operatorname{size}(u, \operatorname{left}) \operatorname{size}(u, \operatorname{right})| 1$

If we call rebuild(u) during an insertion, it is because u is a scapegoat. Assume that path goes to the left:

$$\frac{\text{size}(u.\,\text{left})}{\text{size}(u)} > \frac{2}{3}$$

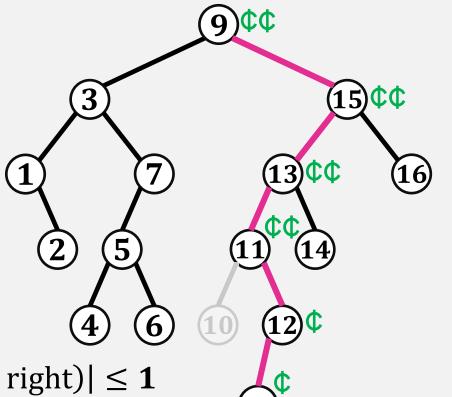
$$size(u) = 1 + size(u.left) + size(u.right)$$

$$size(u. left) - size(u. right) > \frac{1}{3} size(u)$$

The last time a subtree containing u was rebuilt (or when u was inserted), we had  $|\operatorname{size}(u, \operatorname{left}) - \operatorname{size}(u, \operatorname{right})| \leq 1$ 

The number of add/remove operations that have affected u's subtrees since then is at least 1

since then is at least 
$$\frac{1}{3}$$
 size $(u) - 1 \le \#$  credits stored at  $u$ 



• During an insertion or deletion, we give one credit to each node on the path to the inserted (or deleted) node u.

So we hand out at most  $\log_{3/2} q \le \log_{3/2} m$  credits per operation.

ullet During a deletion we also store an additional credit with  $oldsymbol{q}$  .

Thus, in total we give out at most  $O(m \log m)$  credits.

#### Theorem 8.1

A **ScapegoatTree** implements the **SSet** interface. Ignoring the cost of rebuild(u) operations, a **ScapegoatTree** supports the operations  $\operatorname{add}(x)$ , remove(x), and find(x) in  $O(\log n)$  time per operation. Furthermore, beginning with an empty **ScapegoatTree**, any sequence of m add(x) and remove(x) operations results in a total of  $O(m \log m)$  time spent during all calls to rebuild(u).