



Red-Black Trees

Alina Shaikhet

Red-Black Trees

Red-Black Trees are the underlying implementation of **TreeSet** and **TreeMap**

- **Red-Black Trees** are a variation of **binary search trees** with logarithmic height.
- They are one of the most widely used data structures.
 1. A red-black tree storing n values has height at most $2 \log n$.
 2. The $\text{add}(x)$ and $\text{remove}(x)$ operations on a red-black tree run in $O(\log n)$ **worst-case** time.
 3. The amortized number of rotations performed during an $\text{add}(x)$ or $\text{remove}(x)$ operation is **constant**.


Recall SSet implementations:

- **SkipLists** and **Treaps** rely on randomization – $O(\log n)$ running times are only **expected**.
- **ScapegoatTrees** have a guaranteed bound on their height but $\text{add}(x)$ and $\text{remove}(x)$ operations only run in $O(\log n)$ **amortized** time.

2-4 Trees

A **2-4 tree** is a rooted tree with the following properties

- All leaves have the same depth
- Every internal (non-leaf) node has 2, 3, or 4 children.



Claim: If a **2-4 tree** has height h then it has at least 2^h leaves.

Claim: If a **2-4 tree** has $n + 1$ leaves then it has height at most $\log_2(n + 1)$.

$$\# \text{ leaves} \geq 2^h$$

$$n + 1 \geq 2^h$$
$$\log_2(n + 1) \geq h$$


2-4 Trees – find(x)

$O(\log n)$

Internal nodes hold the keys values.

External nodes (leaves) do not store any keys.


keys = **21**
leaves = **22**



Claim: If a **2-4 tree** has n keys then it has exactly $n + 1$ leaves.

2-4 Trees – add(x)


add(68)



2-4 Trees – add(x)

add(68)

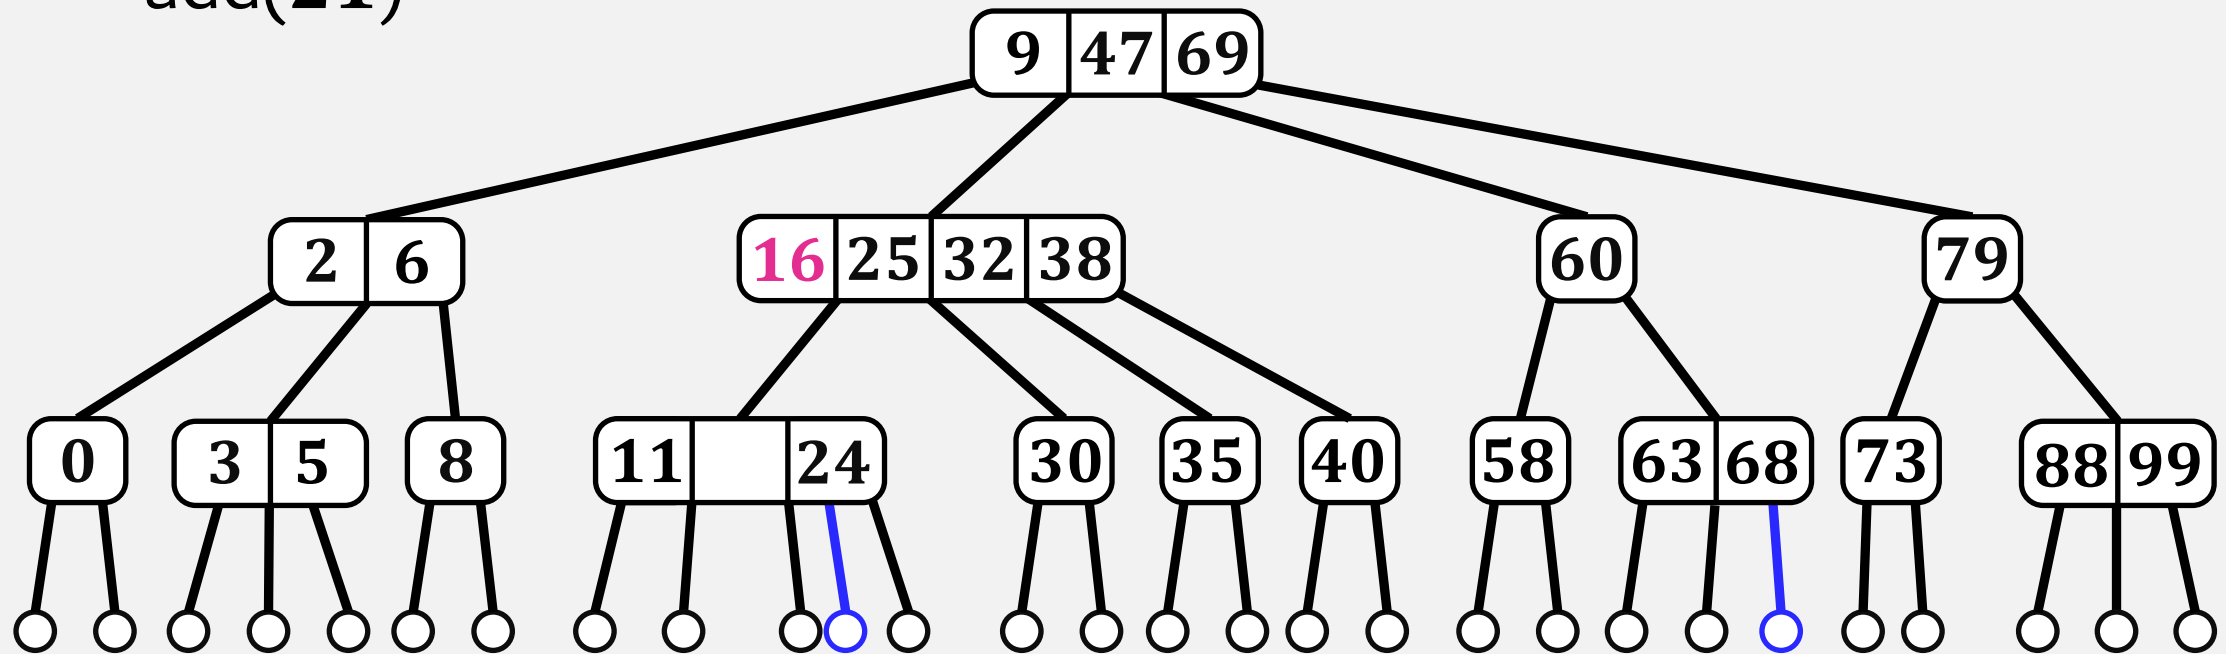
add(21)



2-4 Trees – add(x)

add(68)

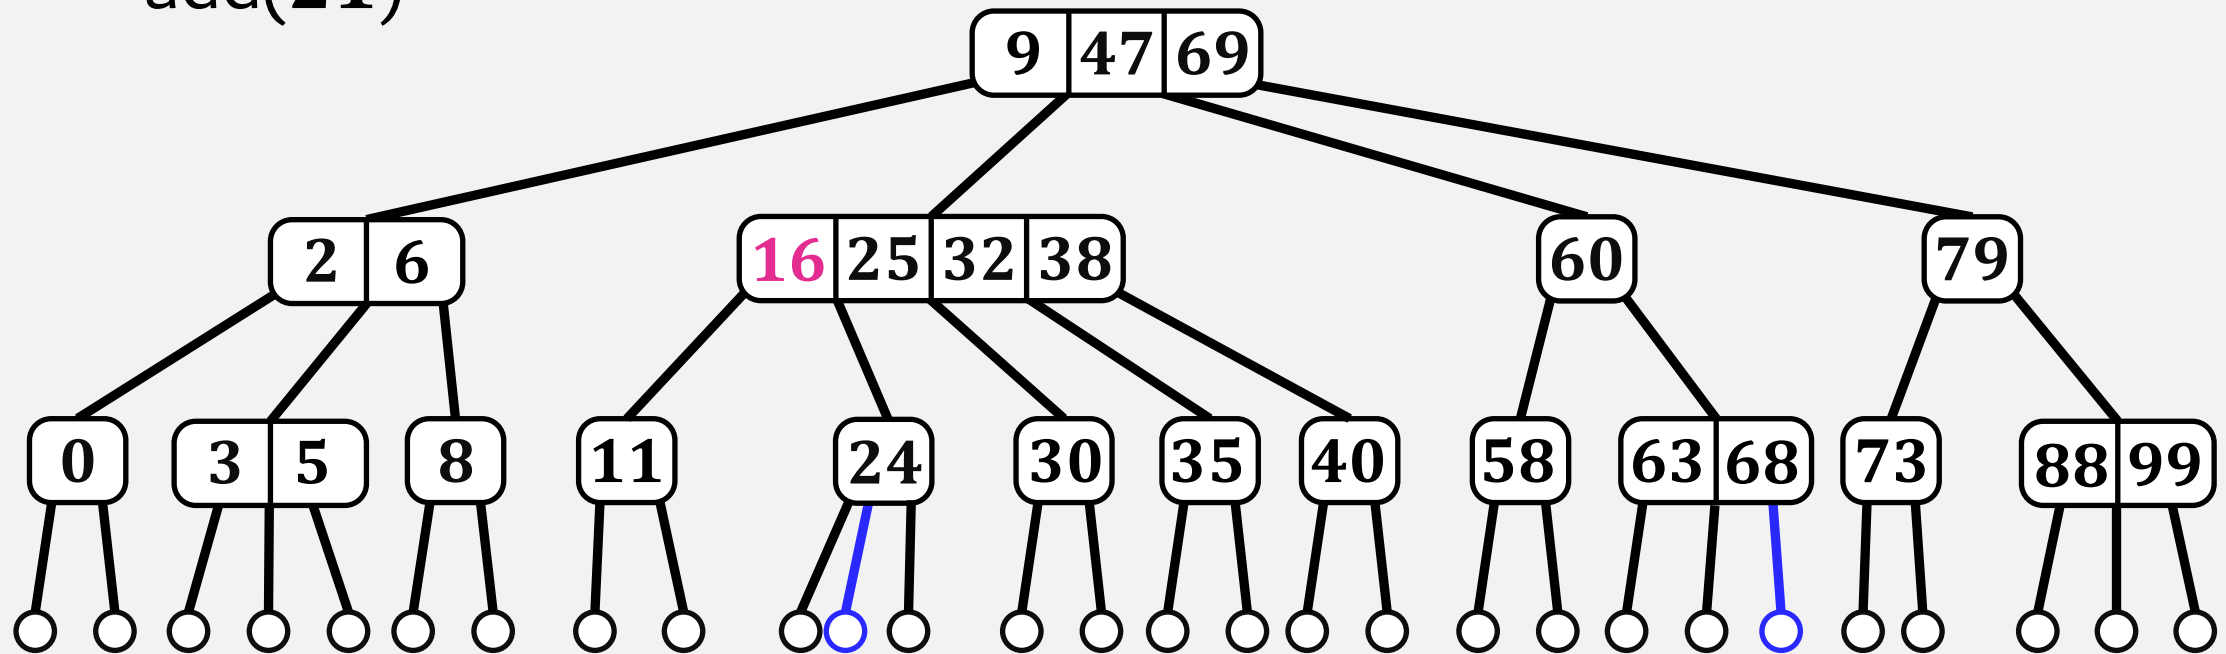
add(21)



2-4 Trees – add(x)

add(68)

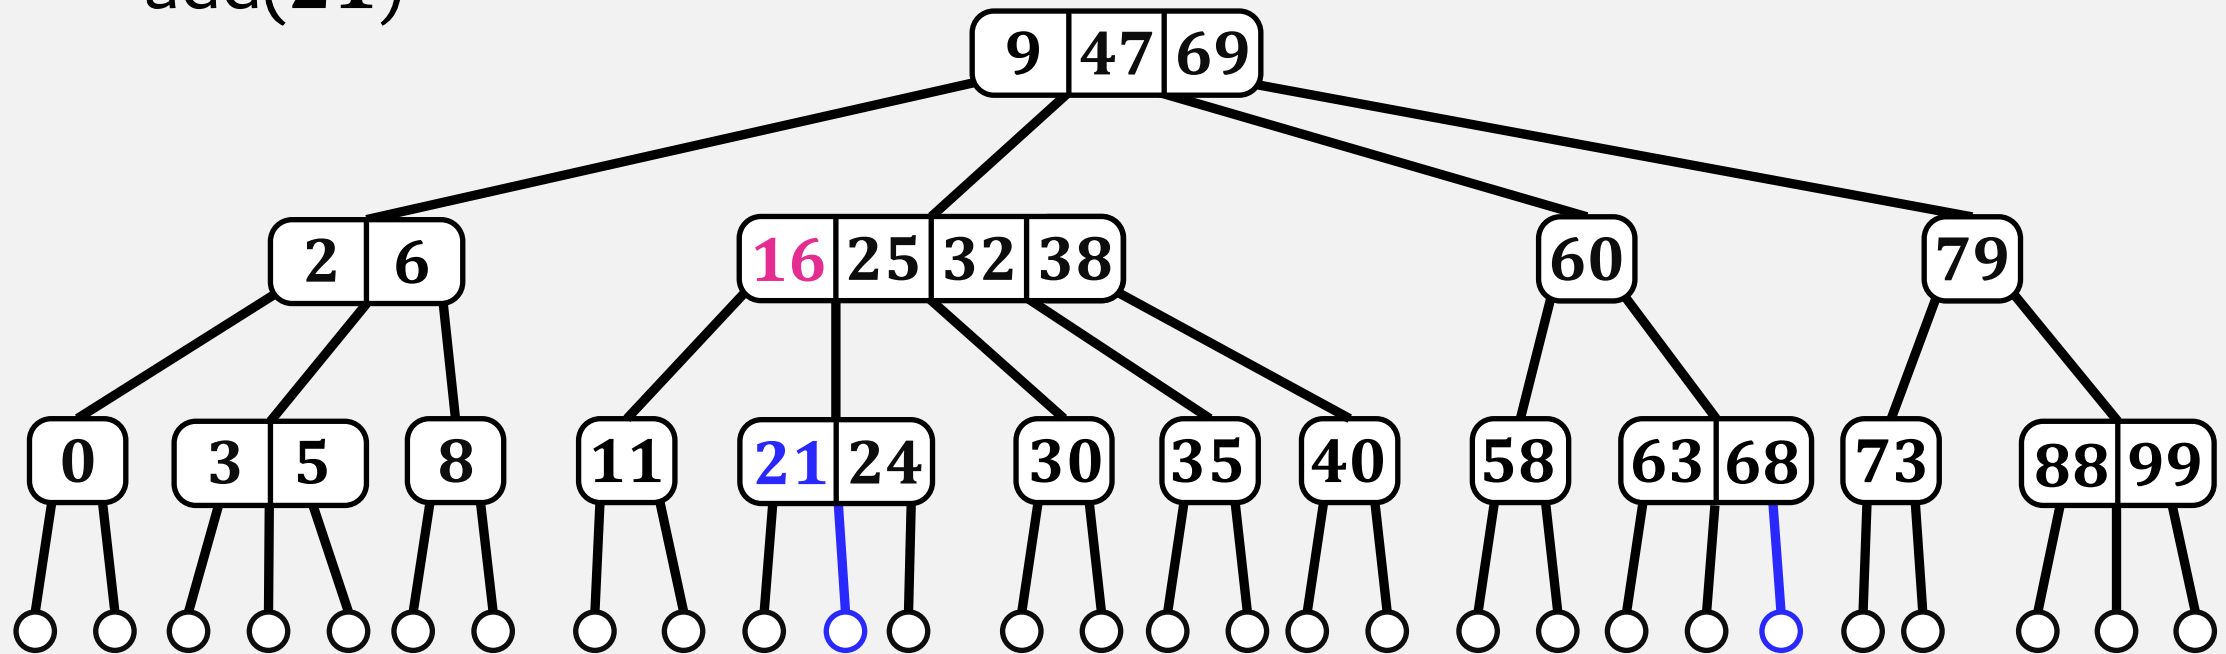
add(21)



2-4 Trees – add(x)

add(68)


add(21)



2-4 Trees – add(x)

add(68)


add(21)



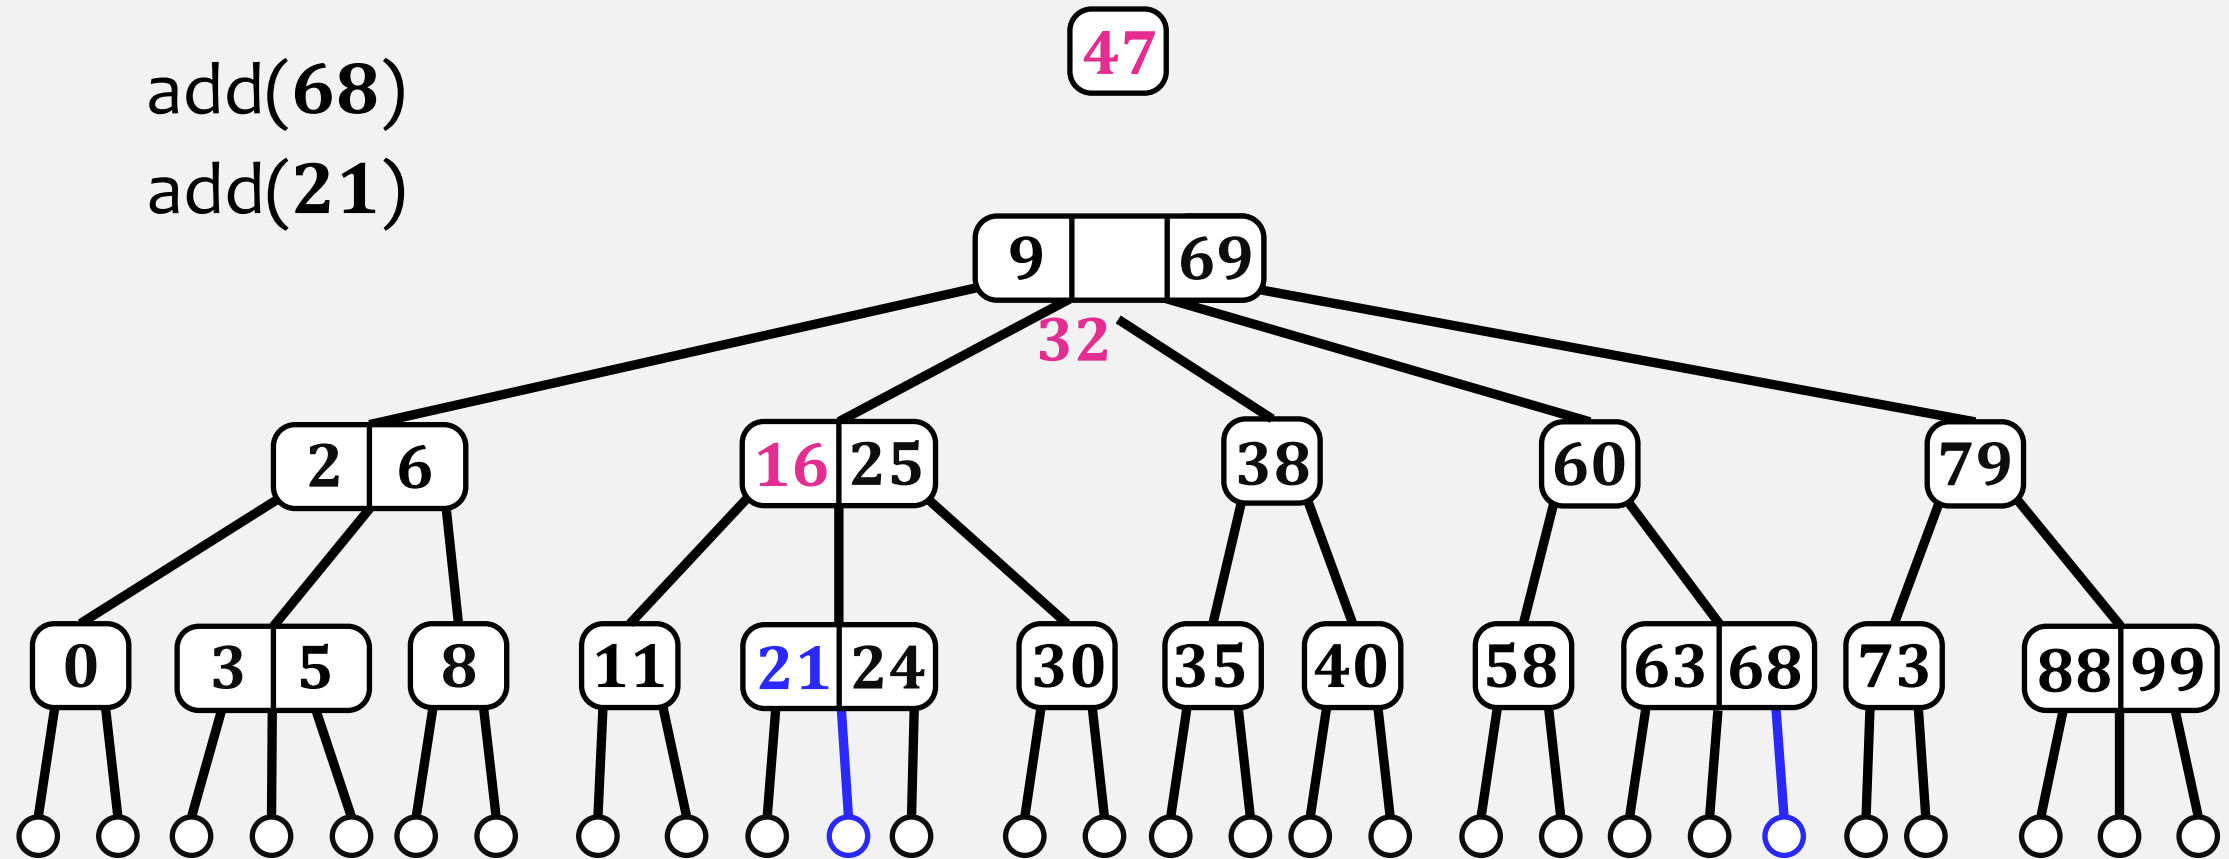
2-4 Trees – add(x)

add(68)

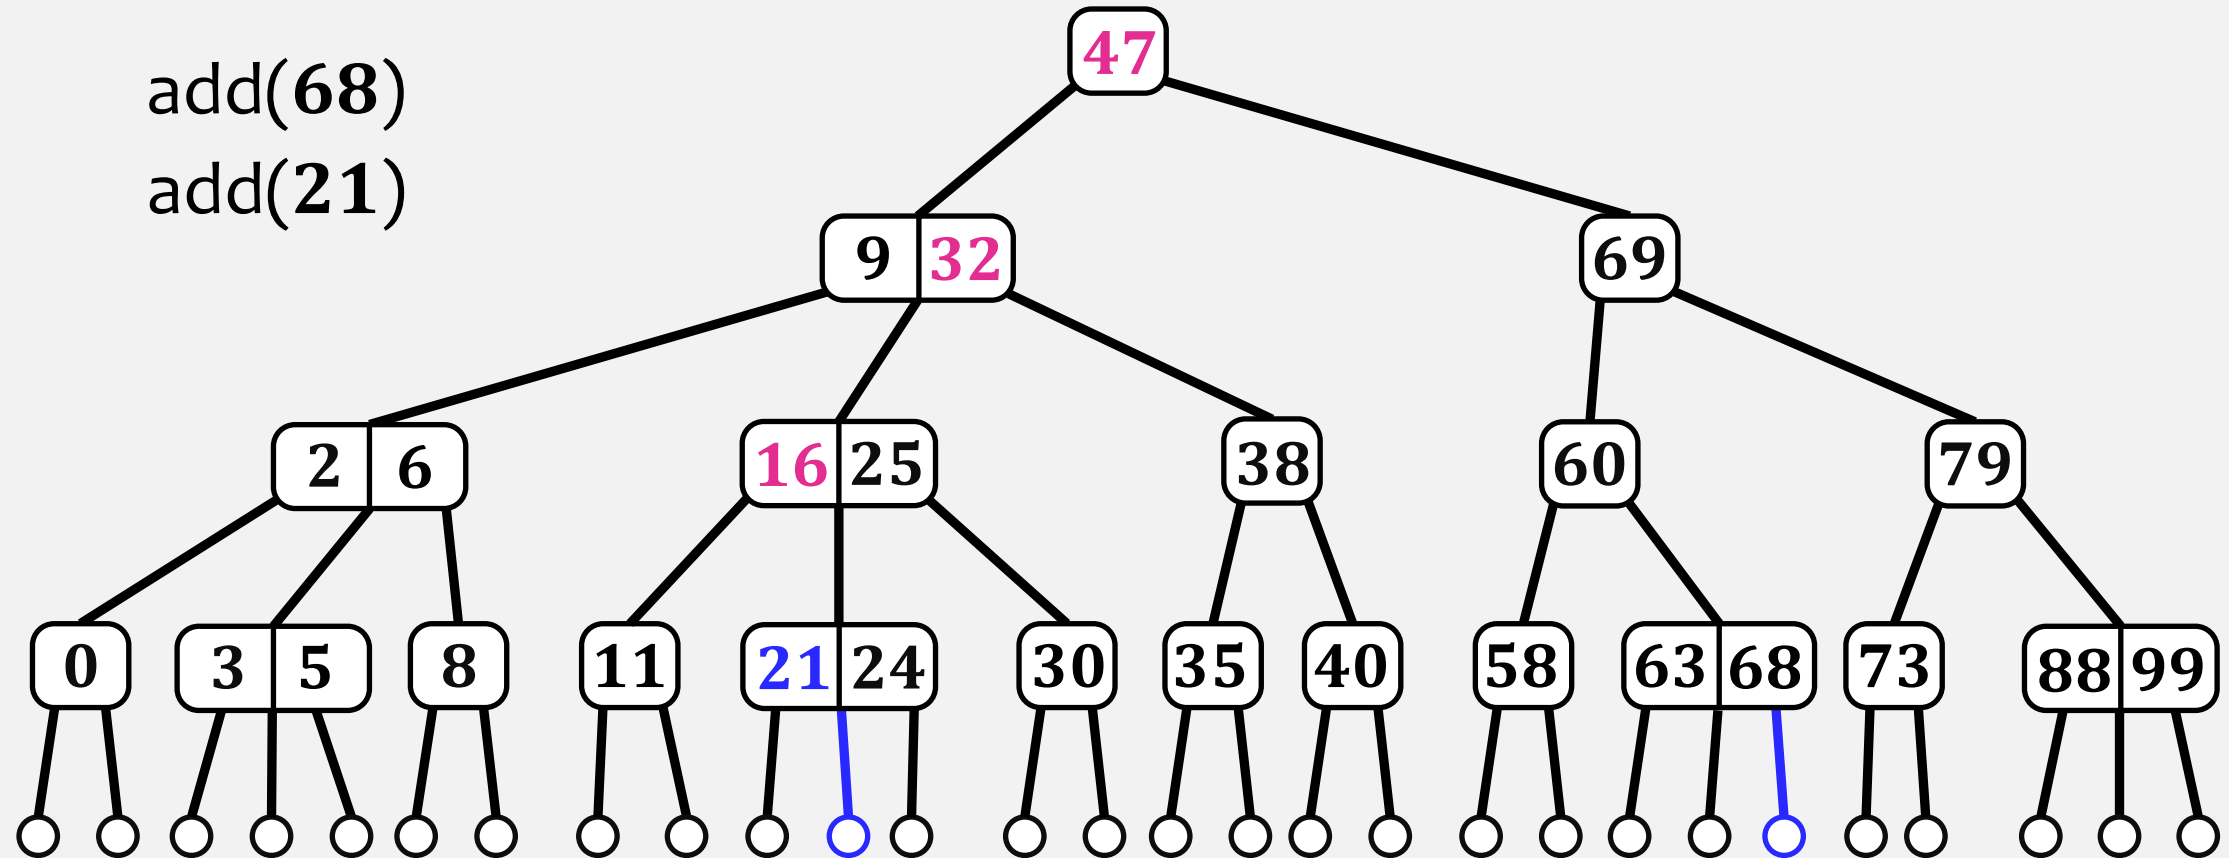
add(21)



2-4 Trees – add(x)



2-4 Trees – add(x)



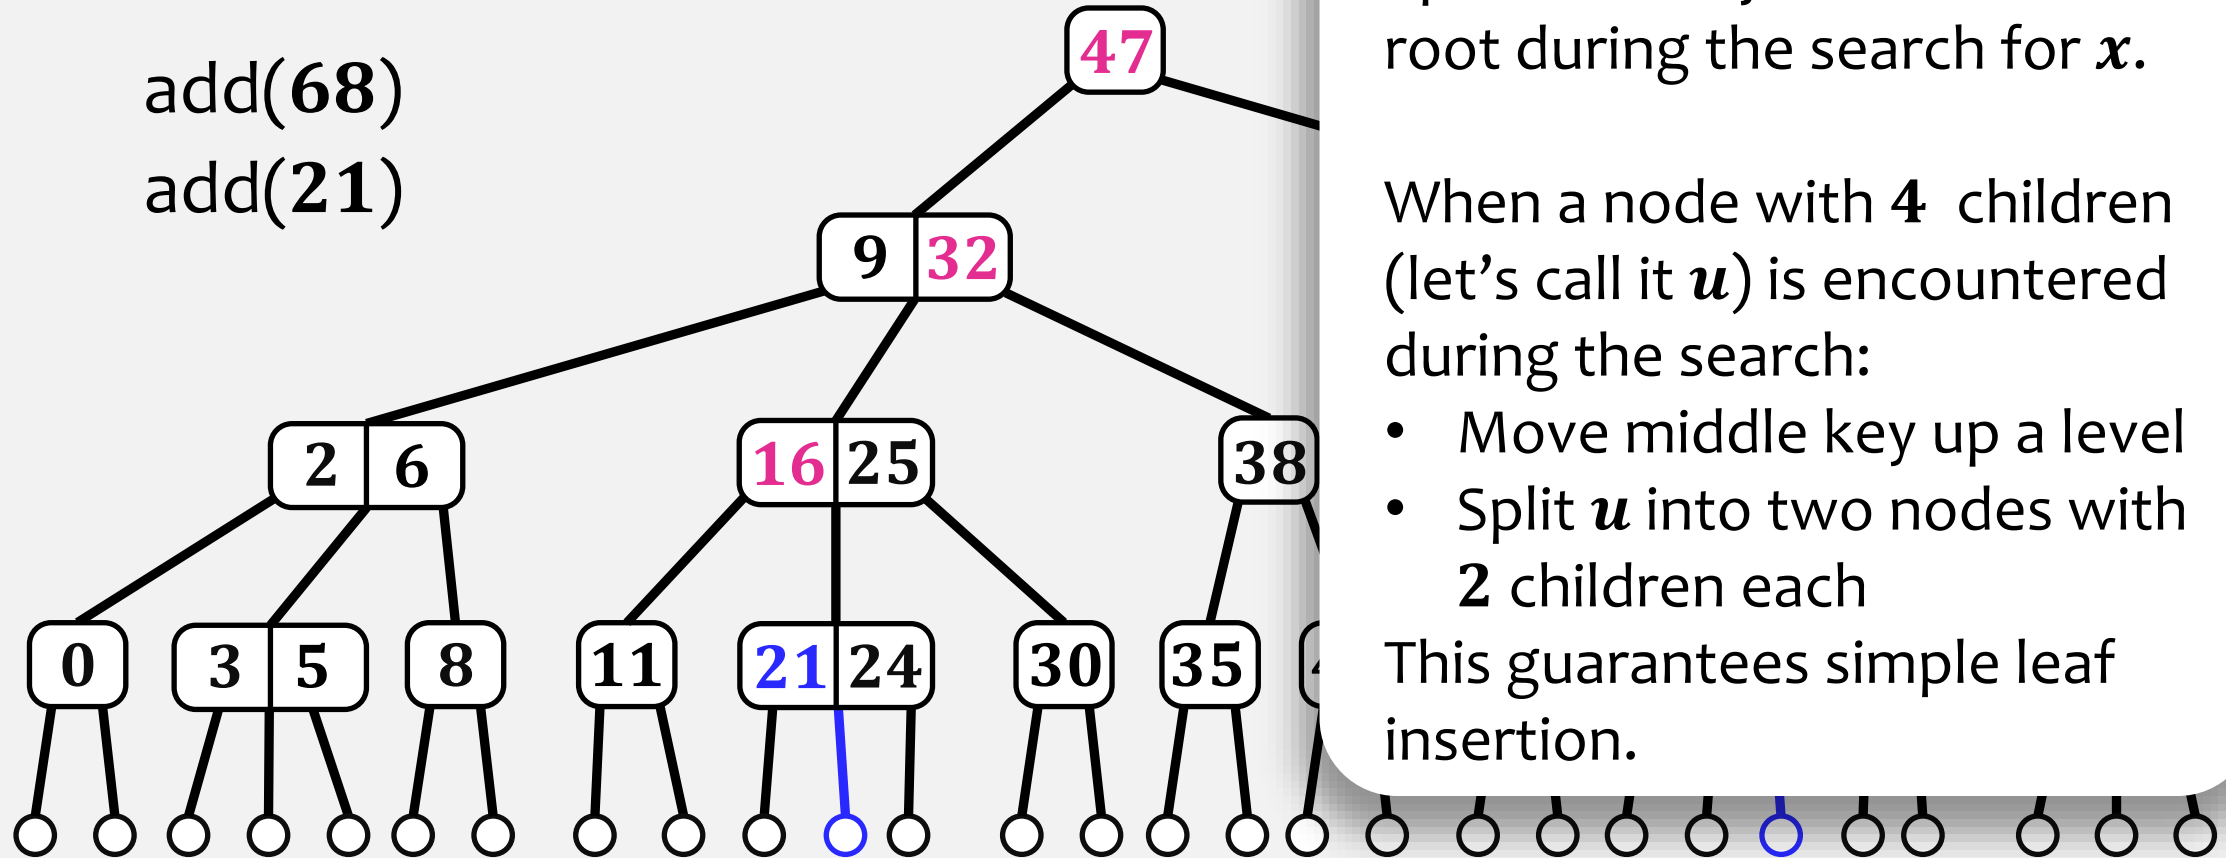
Insertion can be done in one pass

$O(\log n)$

2-4 Trees – add(x)

add(68)

add(21)



Implementation:

Nodes with 4 children are split up on the way down from the root during the search for x .

When a node with 4 children (let's call it u) is encountered during the search:

- Move middle key up a level
- Split u into two nodes with 2 children each


This guarantees simple leaf insertion.

Insertion can be done in one pass

$O(\log n)$


2-4 Trees – remove(x)

remove(8)




2-4 Trees – remove(x)

remove(8)



2-4 Trees – remove(x)


remove(8)



2-4 Trees – remove(x)

To implement remove(x) we need to know:

1. how to borrow a child from a sibling (that has two or three keys)
2. how to merge two nodes (their parent can become empty)



2-4 Trees – remove(x)

Deletion can be tricky.

Removing a node with **2** children involves **rotations** and/or **fusions** (recombine two nodes with **2** children each back into a node with **4** children)

Deletion can be done in a single pass

$O(\log n)$

In many cases instead of deleting consider leaving in place, just mark as deleted. You may reuse these nodes/keys in future insertions.

This approach is not good when doing many add/remove ops of **different** values.

RedBlackTree

red = 0
black = 1

A **red-black tree** is a **binary search tree** such that each node u has a colour which is either **red** or **black**.

Properties:


1. **Black-height:** There are the same number of **black** nodes on every root-to-leaf path.

The sum of the colours on any root-to-leaf path is the same.

2. **No-red-edge:** No two **red** nodes are adjacent.
Each **red** node has a parent, that is a **black** node.

The root is a **black** node.

For any node u , except the root,
 $u.\text{colour} + u.\text{parent.colour} \geq 1$.



Internal nodes hold the keys values.
External nodes (leaves) do not store any keys.

RedBlackTree


Transformation:

Remove each **red** node u and connect its two children directly to its parent.



For each **red** node:

- Parent node is **black**
- Both children nodes are **black**



After the transformation there will be no **red** nodes in the tree.





RedBlackTree



RedBlackTree




RedBlackTree




RedBlackTree

RedBlackTree simulates 2-4 tree using binary tree



2-4 tree




- All leaves have the same depth because of the **black-height** property,
- Every internal (non-leaf) node has **2, 3, or 4** children.

Height

A **2-4 tree** has $n + 1$ leaves that correspond to the $n + 1$ external nodes of the red-black tree.

We know that **2-4 tree** has height at most $\log_2(n + 1)$.



Every root-to-leaf path in the **2-4 tree** corresponds to a path from the root of the **red-black tree** to an external node.



- first and last node in this path are **black** nodes
- at most one (out of every two) internal nodes is a **red** node

So, this path has at most $\log_2(n + 1)$ **black** nodes,
and at most $\log_2(n + 1) - 1$ **red** nodes.

Thus, the height of the red-black tree is at most $2\log_2(n + 1) - 2 \leq 2\log_2 n$

Red-Black Trees – add(x)

add(9)





In a **red-black tree** we need a method of simulating splitting a node with five children in a **2-4 tree**:

We can “split” the node by recolouring: node 7 becomes red; nodes 6, 8 become black

Red-Black Trees – add(x)

add(9)





In a **red-black tree** we need a method of simulating splitting a node with five children in a **2-4 tree**:

We can “split” the node by recolouring: node 7 becomes red; nodes 6, 8 become black

Red-Black Trees – add(x)

add(9)





In a **red-black tree** we need a method of simulating splitting a node with five children in a **2-4 tree**:

We can “split” the node by recolouring: node 7 becomes red; nodes 6, 8 become black

Red-Black Trees – add(x)

add(9)



In a **red-black tree** we need a method of simulating splitting a node with five children in a **2-4 tree**:

We can “split” the node by recolouring: node 7 becomes red; nodes 6, 8 become black

Red-Black Trees – remove(x)

there are many cases that must be considered

To implement remove(x) we need to know:

1. how to merge two nodes and

Merging two nodes is the inverse of a split. We can do it by colouring two (**black**) siblings **red** and colouring their (**red**) parent **black**.

2. how to borrow a child from a sibling

Borrowing from a sibling is the most complicated of the procedures and involves both **rotations** and **recoloring** nodes.

We also need to maintain **no-red-edge** property and the **black-height** property.

Theorem 9.1

A **RedBlackTree** implements the **SSet** interface and supports the operations $\text{add}(x)$, $\text{remove}(x)$, and $\text{find}(x)$ in $\mathcal{O}(\log n)$ worst-case time per operation.